

REGULARIZED REGRESSION

Kwabena Asabere

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
from sklearn.linear_model import Ridge,Lasso,RidgeCV,LassoCV,ElasticNetCV,ElasticNet
from sklearn.preprocessing import StandardScaler,OneHotEncoder,PolynomialFeatures,FunctionTransformer
from sklearn.pipeline import Pipeline,make_pipeline
from sklearn.model_selection import train_test_split,cross_val_score,cross_validate,cross_val_score
from sklearn.metrics import root_mean_squared_error,mean_squared_error
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
```

```
housing = pd.read_csv(r"C:\Users\KAsab\Desktop\PYTHON\california_housing.csv",engine = "python")
```

```
housing.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_value
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	83200.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	84500.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	71600.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	51300.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	31600.0

```
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
```

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	longitude	20640 non-null	float64
1	latitude	20640 non-null	float64
2	housing_median_age	20640 non-null	float64
3	total_rooms	20640 non-null	float64
4	total_bedrooms	20433 non-null	float64
5	population	20640 non-null	float64
6	households	20640 non-null	float64
7	median_income	20640 non-null	float64
8	median_house_value	20640 non-null	float64
9	ocean_proximity	20640 non-null	object

dtypes: float64(9), object(1)

memory usage: 1.6+ MB

```
housing.isna().sum()
```

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	207
population	0
households	0
median_income	0
median_house_value	0
ocean_proximity	0

dtype: int64

```
housing["total_bedrooms"]=housing["total_bedrooms"].fillna(housing["total_bedrooms"].median())
```

```
housing.isna().sum()
```

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	0
population	0
households	0

```

median_income      0
median_house_value  0
ocean_proximity    0
dtype: int64

```

```

housing["rooms_per_house"] = housing['total_bedrooms']/housing["total_rooms"]
housing["people_per_house"] = housing["population"]/housing["households"]

```

```
housing.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3254
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.0699
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.1401
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6414
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.7861

```

numeric_features = ["rooms_per_house","people_per_house"]
log_features = housing.iloc[:,8].columns.to_list()
cat_features = housing.select_dtypes(include = "object").columns.to_list()

```

```

X = housing.drop(columns = ["median_house_value"])
y = housing["median_house_value"]

```

```
X_train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 16512 entries, 14196 to 15795
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
---  -
0   longitude            16512 non-null  float64
1   latitude             16512 non-null  float64
2   housing_median_age   16512 non-null  float64
3   total_rooms          16512 non-null  float64
4   total_bedrooms       16512 non-null  float64
5   population           16512 non-null  float64
6   households            16512 non-null  float64
7   median_income        16512 non-null  float64

```

```

8   ocean_proximity      16512 non-null  object
9   rooms_per_house      16512 non-null  float64
10  people_per_house     16512 non-null  float64
dtypes: float64(10), object(1)
memory usage: 1.5+ MB

```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```

numeric_transformer = Pipeline( steps = [
    ("imputer",SimpleImputer(strategy = "mean")),
    ("scaler",StandardScaler())
])

log_transformer = Pipeline(steps = [
    ("imputer",SimpleImputer(strategy = "mean")),
    ("log",FunctionTransformer(np.log1p))
])

cat_transformer = Pipeline(steps = [
    ("imputer",SimpleImputer(strategy = "most_frequent")),
    ("onehot",OneHotEncoder(handle_unknown = "ignore",sparse_output = False))
])

```

```

preprocessor = ColumnTransformer(transformers =[
    ("log",log_transformer,log_features),
    ("num",numeric_transformer,numeric_features),
    ("cat",cat_transformer,cat_features)
])

```

```

pipeline = Pipeline( steps = [
    ("preprocessor",preprocessor),
    ("ridge",Ridge())
])

```

```

param_grid = {
    "ridge__alpha": [0.1,1.0,10,100,1000]
}

```

The names of the hyperparameters in the `param_grid` must contain the name of the step in the pipeline followed by 2 underscores and then the hyperparameter.